


(a) finer LSM-D helps only under convex cost



(b) RL-kernel tracks even the $M = 17$ LSM; no-kernel sits below

